

Universal Semi-Supervised Semantic Segmentation

A novel approach towards learning efficient and transferable representations
for dense prediction tasks

Tarun Kalluri² Girish Varma¹ CV Jawahar¹
CVIT, IIIT Hyderabad¹

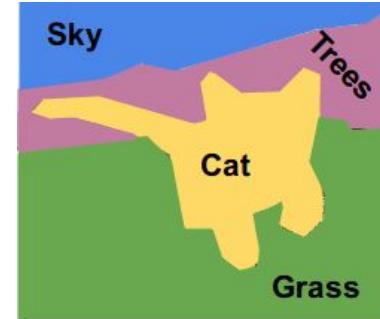
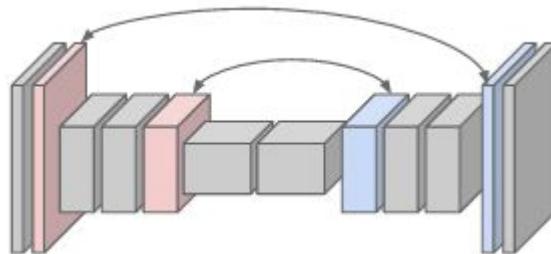
Manmohan Chandraker²
UCSD²

Universal Semi-Supervised Semantic Segmentation



Universal Semi-Supervised Semantic Segmentation

- Semantic segmentation involves pixel-level semantic labeling of images.
- Applications include autonomous driving, remote sensing, satellite imagery and medical imaging.
- Dense predictions require dense annotations.
 - Challenging and tedious task to label.
 - Requires massive human effort.
 - Unlabeled data easier to obtain.



Universal **Semi-Supervised** Semantic Segmentation

- Semi-Supervised learning involves learning with few labeled training data and lots of unlabeled data.
- Unlabeled data assist the labeled data in learning the classifiers.
- Many techniques to do this.
 - Pseudo-Labeling
 - Co-Training
 - Entropy Minimization

Universal **Semi-Supervised** Semantic Segmentation

- Semi-Supervised learning involves learning with few labeled training data and lots of unlabeled data.
- Unlabeled data assist the labeled data in learning the classifiers.
- Many techniques to do this.
 - Pseudo-Labeling - Assign labels to unlabeled set using model trained on labeled set.
 - Co-Training
 - Entropy Minimization

Universal **Semi-Supervised** Semantic Segmentation

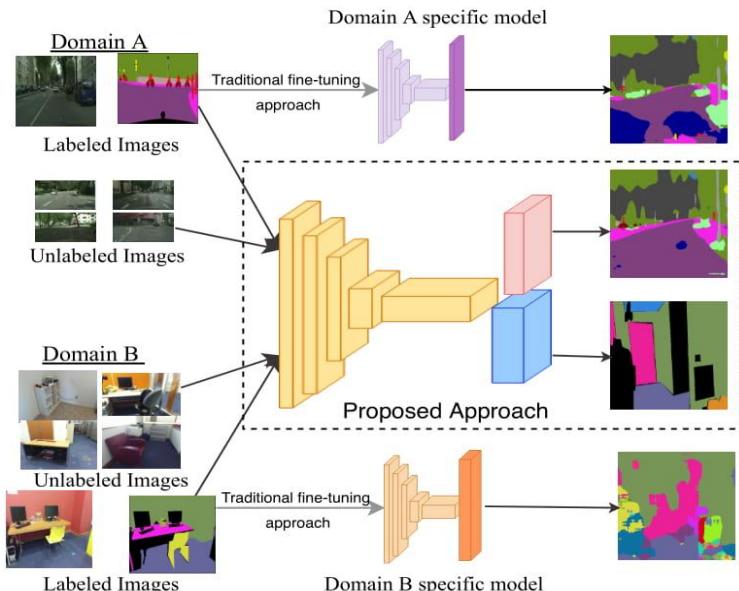
- Semi-Supervised learning involves learning with few labeled training data and lots of unlabeled data.
- Unlabeled data assist the labeled data in learning the classifiers.
- Many techniques to do this.
 - Pseudo-Labeling
 - Co-Training - Generate multiple, consistent views of unlabeled data using different models.
 - Entropy Minimization

Universal **Semi-Supervised** Semantic Segmentation

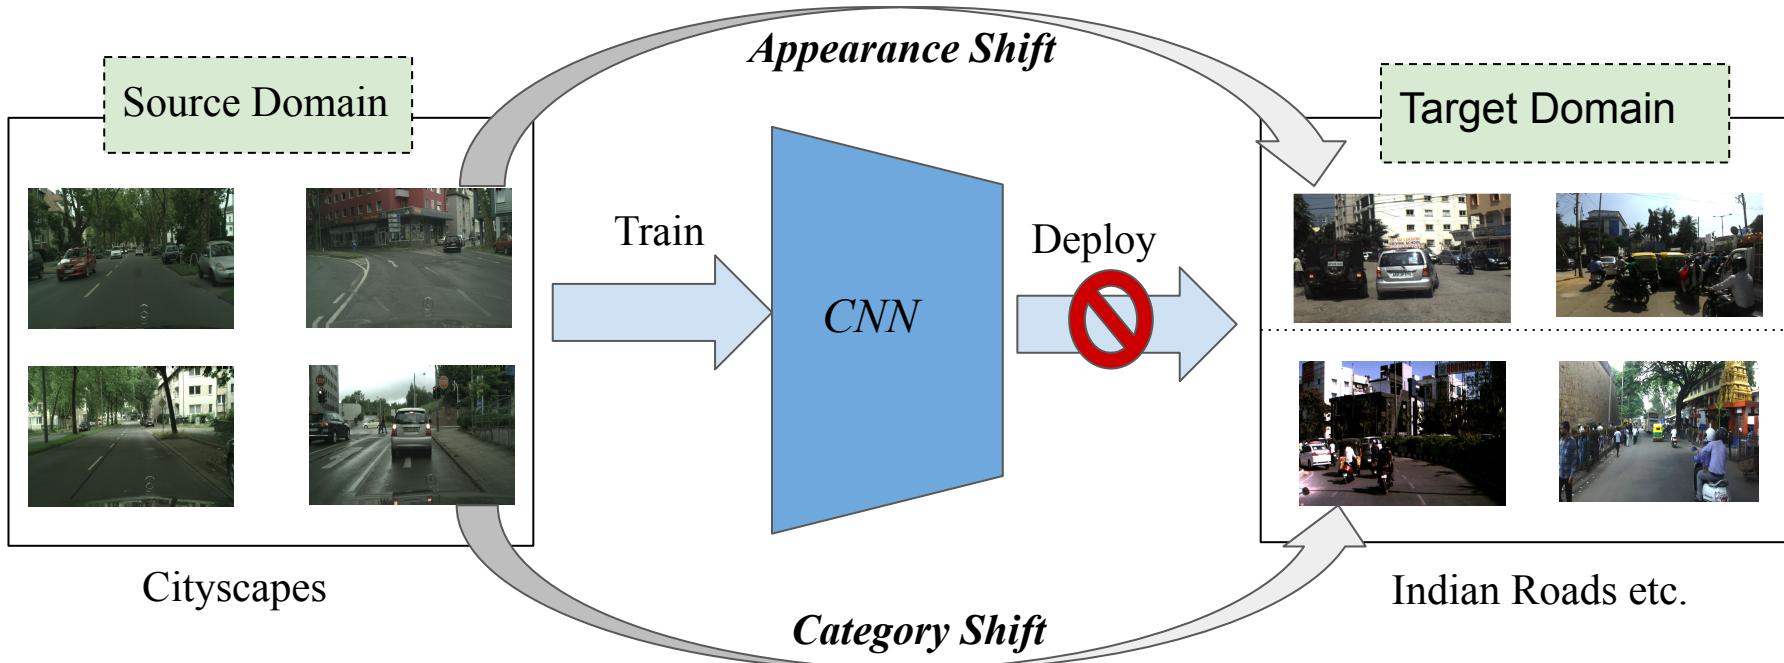
- Semi-Supervised learning involves learning with few labeled training data and lots of unlabeled data.
- Unlabeled data assist the labeled data in learning the classifiers.
- Many techniques to do this.
 - Pseudo-Labeling
 - Co-Training
 - Entropy Minimization - Encourages Low density separation in feature space

Universal Semi-Supervised Semantic Segmentation

- Idea: Obtain a common semantic segmentation model across widely disparate domains having limited labeled data.
- A good universal model should ensure that, across all domains
 - A single model is deployed
 - Unlabeled data is used
 - Performance is improved
 - And label spaces (semantic content) may differ.



Challenge: Domain Shift + Category Shift

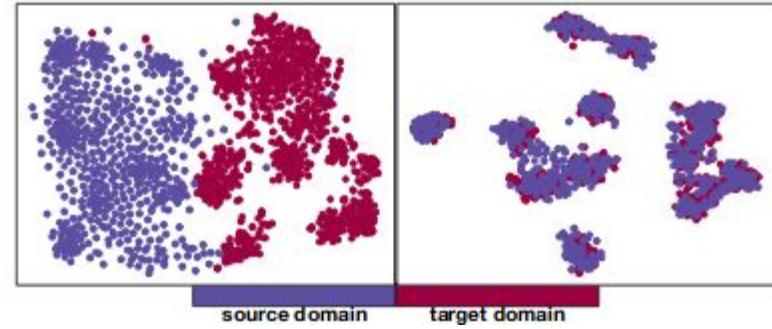


Related Approaches

- Fine-tuning on labeled data.
 - requires separate models for each domain
[catastrophic forgetting].
 - Introduces both training and deployment overhead.

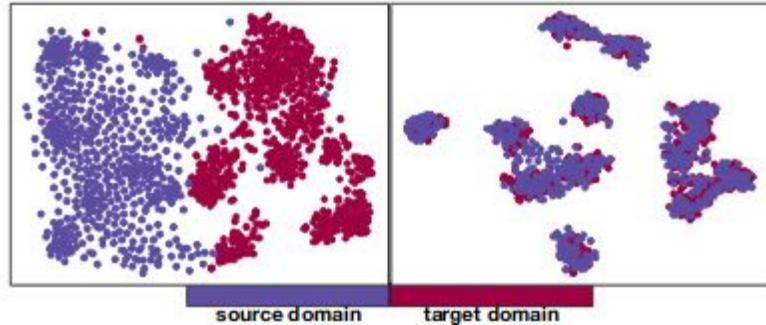
Related Approaches: Domain Adaptation

- Unsupervised domain adaptation tackles the problem of domain shift between a labeled source and unlabeled target domain.
- Different ways include
 - Minimizing MMD between datasets.
 - Aligning features from both the domains through adversarial approaches.



Limitations to Domain Adaptation

- Assumption of same label spaces between the source and target domain.
- Model adapted to one target domain does not apply to other domains.



Other Related Approaches

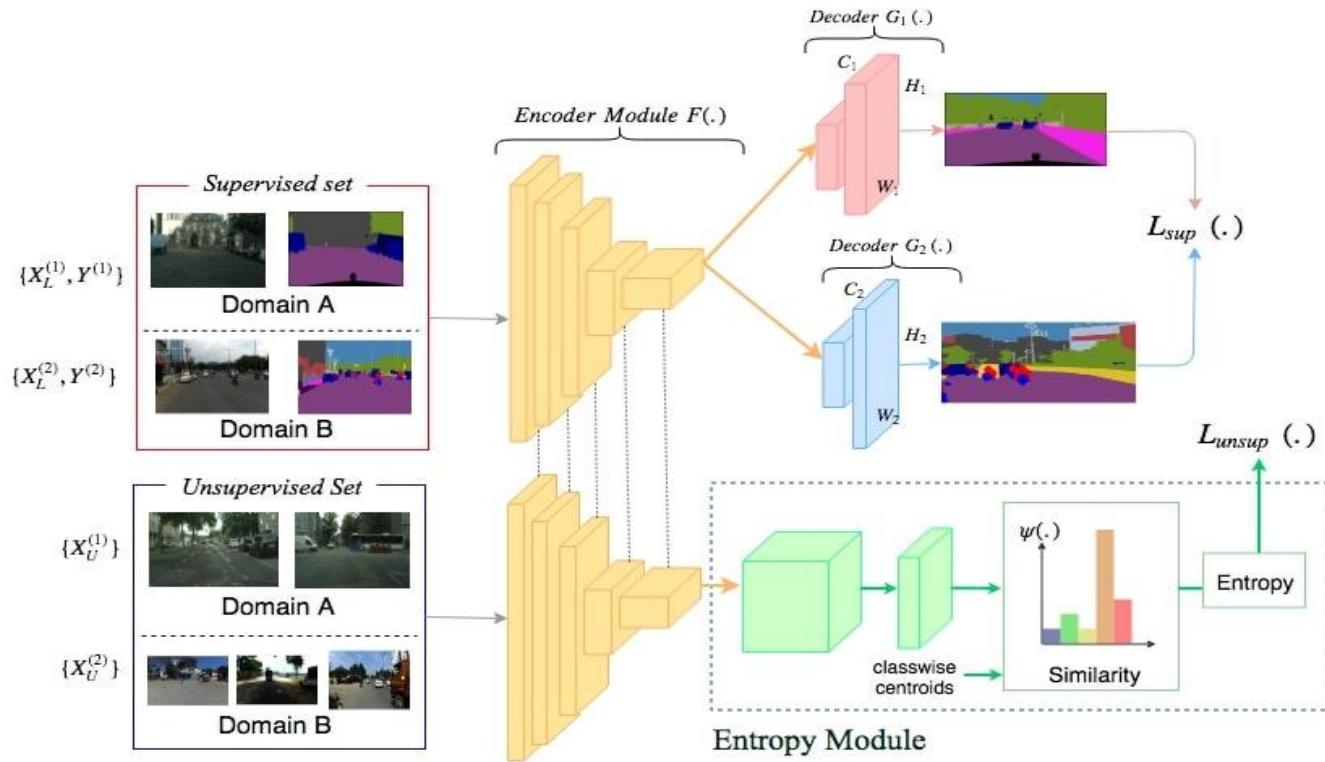
	Source Unlabeled Data	Target Unlabeled Data	Joint Model	Mixed Labels Support
Fine Tuning	✗	✗	✗	✓
Semi-supervised [Hung 2018]	✓	✗	✗	NA
CyCADA [Hoffman 2018]	✗	✓	✓	✗
Joint Training	✗	✗	✓	✓
Our Approach	✓	✓	✓	✓

Prior works fall short in addressing the semantic change, or do not make use of large scale unsupervised images.

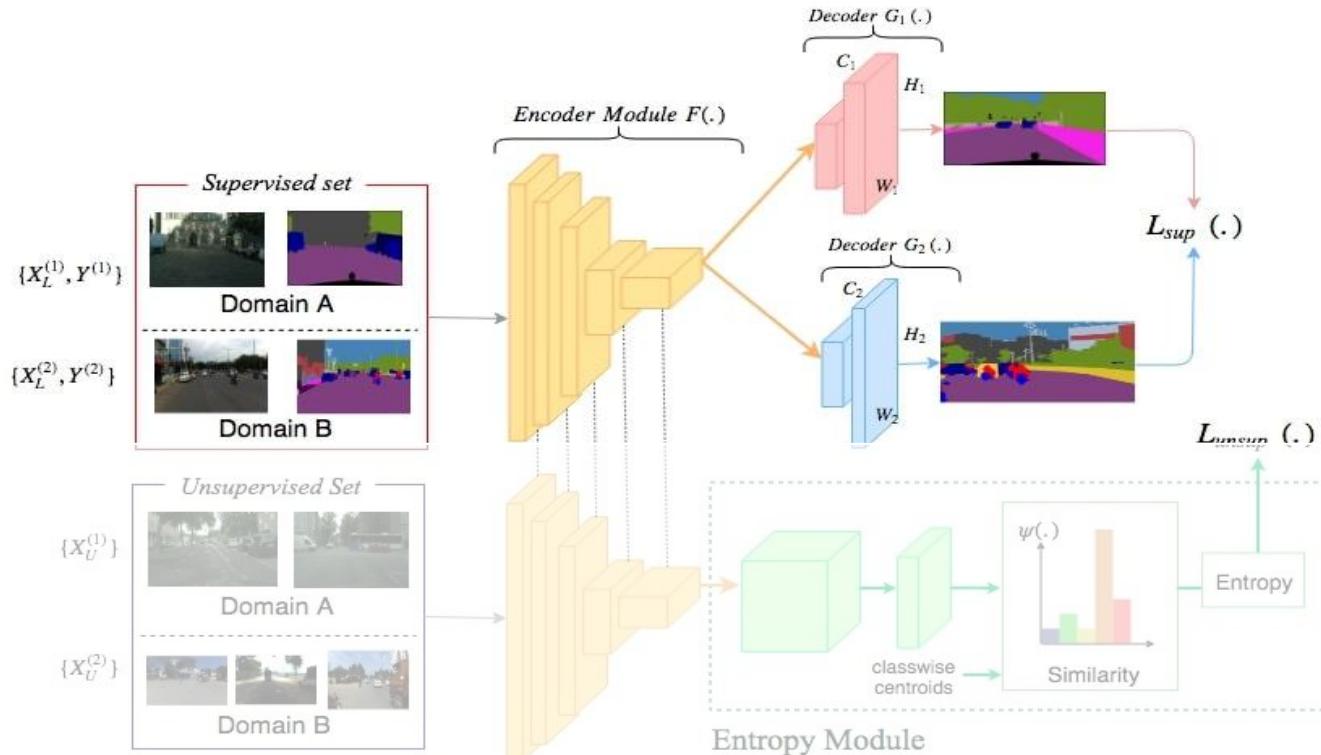
Our Approach

- Use labeled data to build the task specific classifiers.
- Use unlabeled data from each domain to align visually similar features for improved classification.

Proposed Architecture

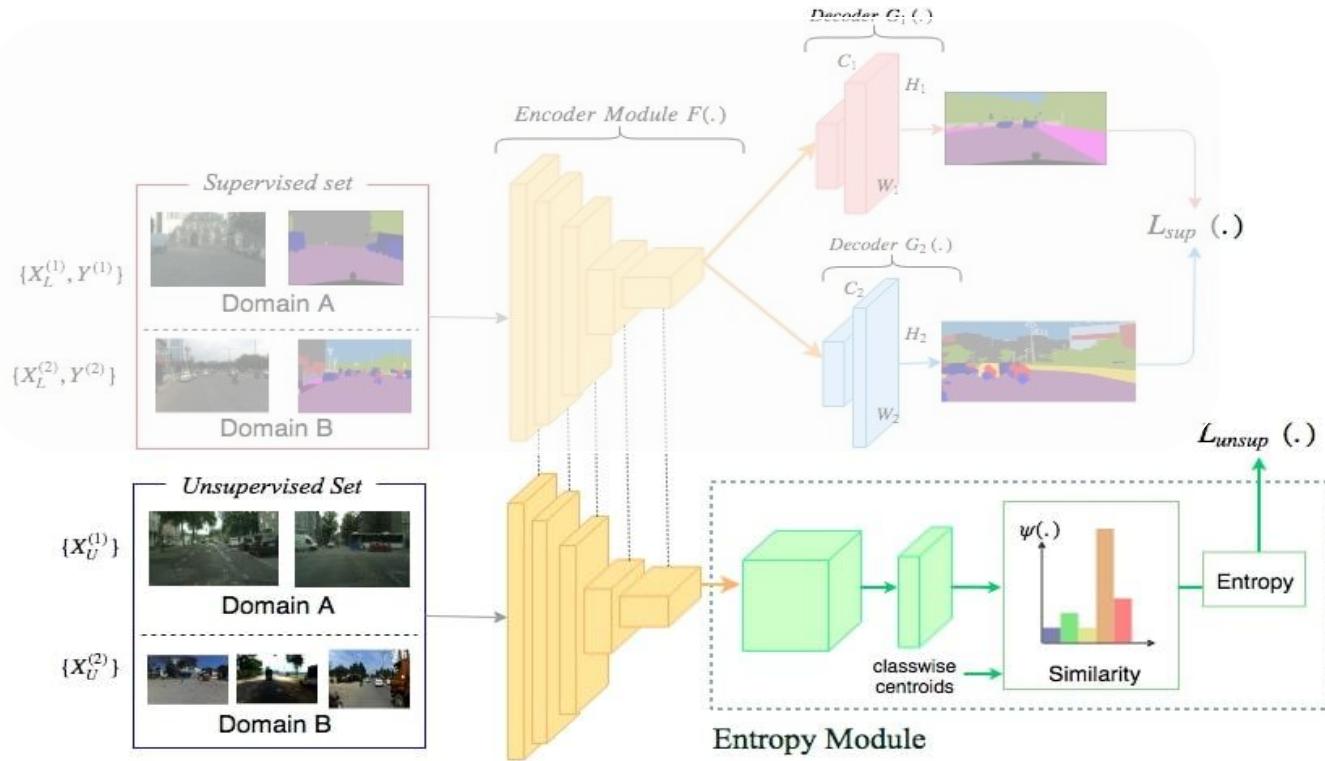


Supervised Loss



$$L_{sup} = \sum_k \frac{1}{N_l^{(k)}} \sum_{x_i \in D^{(k)}} \psi_k(y_i, \mathcal{G}_k(\mathcal{F}(x_i)))$$

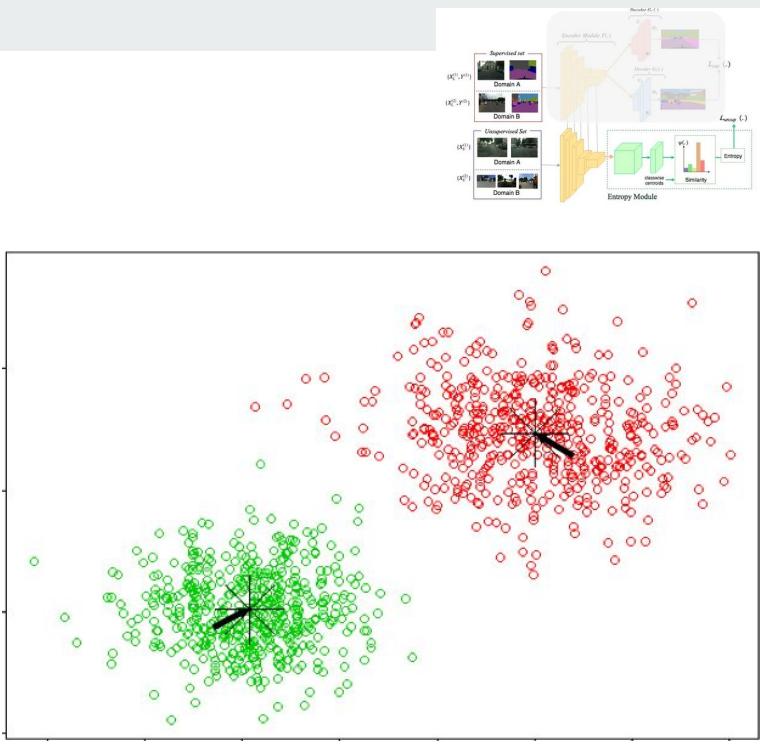
Semi-Supervised Loss



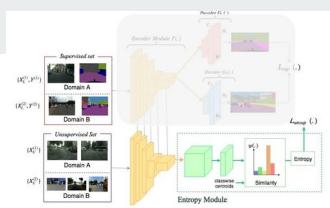
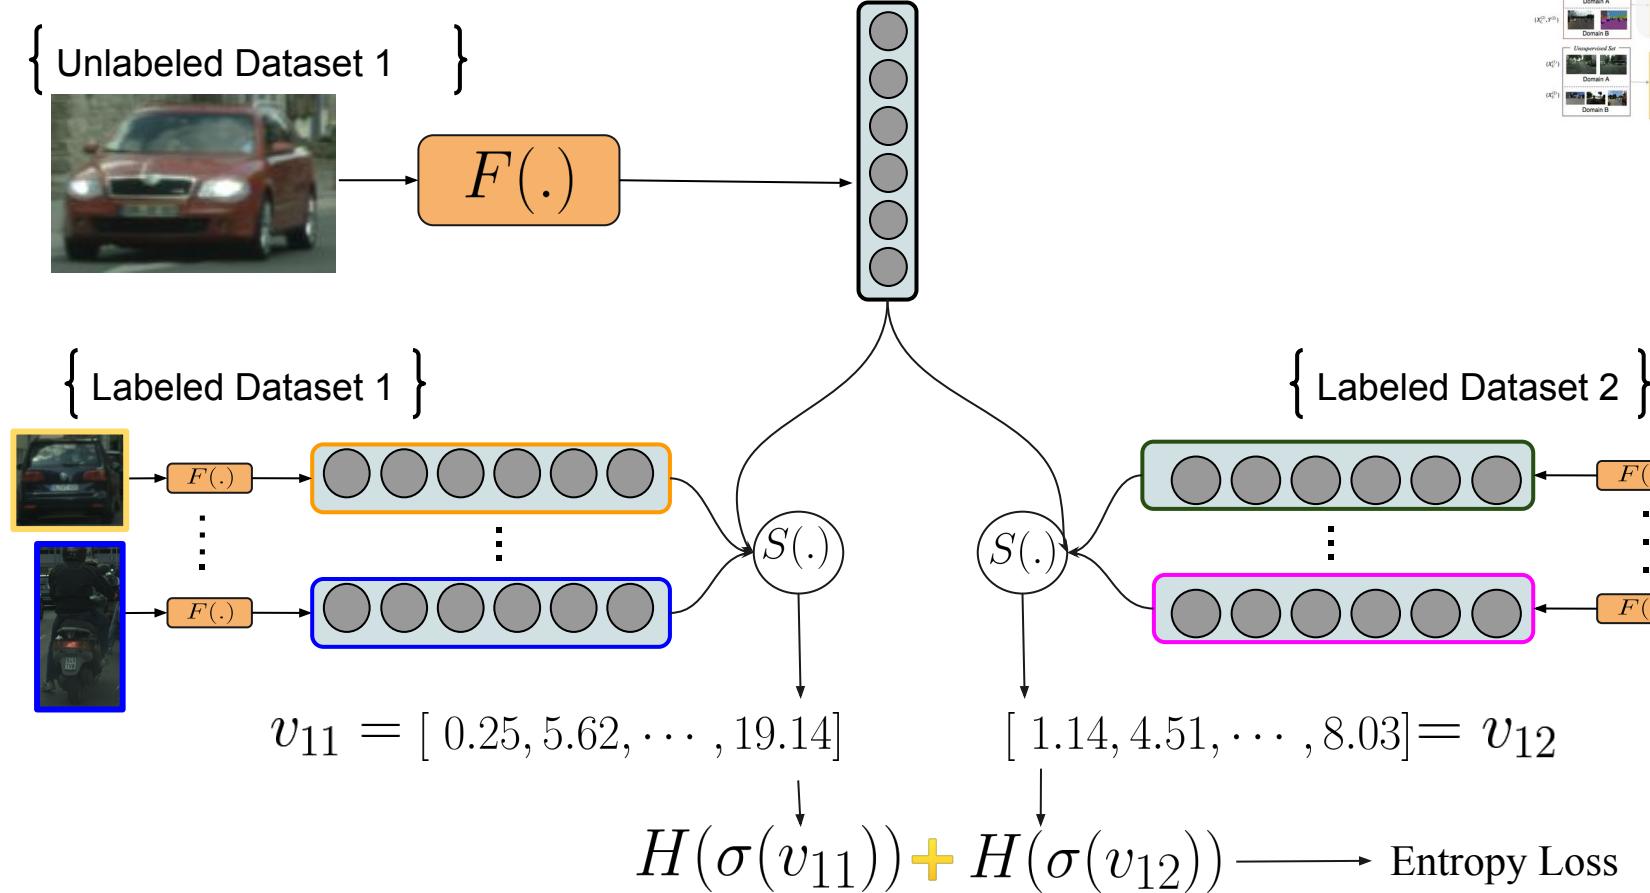
Entropy Module

Entropy Regularization

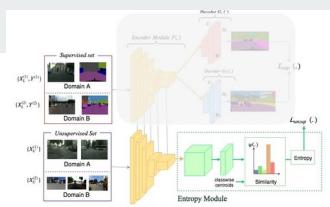
- Entropy regularization (or entropy minimization) encourages low density separation between the clusters in the feature space.
- Results in high confidence predictions and smooth output maps for semi-supervised learning
- We use it to transfer semantic information from unlabeled to labeled images, between both the domains.



Entropy Regularization using similarity matching



Entropy Regularization using similarity matching



- Semantic Segmentation involves feature extraction at pixel level, so entropy matching is done for features computed at each pixel.
- Pixel level similarity with all the images is costly, so we calculate prototype features for each class.
- Dot product is used as the similarity metric $S(\cdot)$
- We perform both within dataset matching and cross dataset matching. This handles disparate label spaces.

Semi-Supervised Entropy Loss

$$L_{u,c} = \mathcal{H}(\sigma([v_{12}]))) + \mathcal{H}(\sigma([v_{21}])))$$

$$L_{u,w} = \mathcal{H}(\sigma([v_{11}]))) + \mathcal{H}(\sigma([v_{22}])))$$

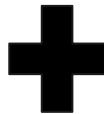
Total Loss / Training Objective

$$L_t = L_S + \alpha \cdot L_{u,c} + \beta \cdot L_{u,w}$$

Datasets for Comparing the Universal Model



Cityscapes



IDD (Indian Roads)



SUN RGB-D (Indoor Scenes)



CamVid



Experimental Results

Method	Road	SideWalk	Building	Wall	Fence	Pole	Traff. lt.	Traff. Sgn.	Veg.	Train	Sky	Person	Rider	Car	Track	Bus	Train	MotorCyc.	Bicycle	mIoU
CS only	91.76	54.78	80.02	3.70	16.58	29.84	22.31	33.74	83.88	32.89	82.07	52.67	21.57	81.11	19.01	3.87	0.0	19.64	49.01	40.97
Univ-basic	87.00	44.54	77.77	10.21	11.07	25.54	14.51	25.82	80.72	22.40	78.19	49.00	19.64	75.35	1.86	0.25	10.98	8.83	41.08	36.04
Univ-full	92.18	51.29	80.07	0.0	24.01	33.73	26.16	38.71	82.30	36.39	81.61	54.38	20.48	81.71	2.37	22.79	3.85	1.31	46.23	41.03

Method	N=375		
	CS	CamVid	Avg.
Train on CS	55.07	48.52	51.80
Train on CVD	26.45	60.61	43.53
Hung <i>et al.</i> [28]	58.80	-	-
Souly <i>et al.</i> [61]	-	58.20	-
Univ-basic (\mathcal{L}_s)	53.14	65.33	59.24
Univ-cross (+ \mathcal{L}_c)	56.36	63.34	59.85
Univ-full (+ $\mathcal{L}_c, \mathcal{L}_w$)	55.92	64.72	60.32

Method	N=100 (Resnet-18)		
	CS	IDD	Avg.
Train on CS	40.97	14.64	27.81
Train on IDD	25.05	26.53	25.79
Univ-basic	37.94	25.21	31.58
Univ-full	36.48	27.45	31.97

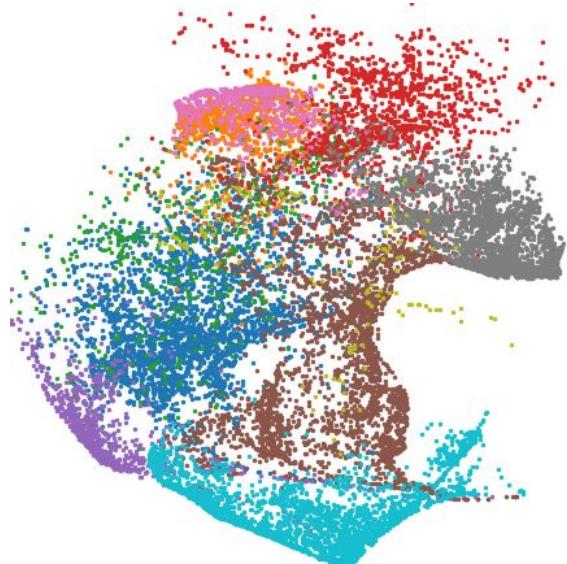


Alignment across environments

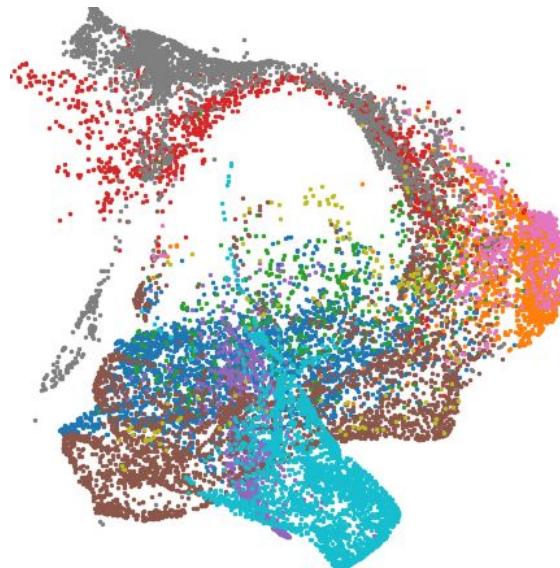
- With only 28% labeled data from SUN dataset and no synthetic examples, we obtain ~88% of performance obtained with full dataset.

Method	Labeled Examples	CS	SUN	Avg.
Train on CS	1.5k	64.23	15.47	39.85
Train on SUN	1.5k	15.61	42.52	29.07
SceneNet [41]	Full(5.3k)	-	49.8	-
Univ-basic	1.5k	58.01	31.55	44.78
Ours[Univ-full]	1.5k	57.91	43.12	50.52

tSNE Embedding Visualization



(a) Without Entropy Module



(b) With Entropy Module

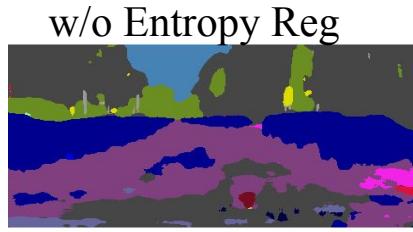


Qualitative Results

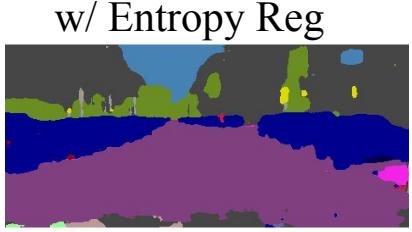
Cityscapes



Input Image



w/o Entropy Reg



w/ Entropy Reg

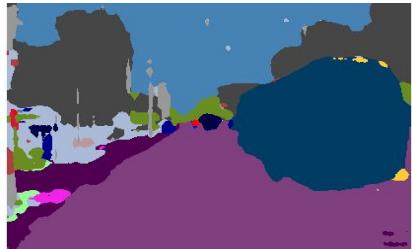


Ground Truth

IDD



w/o Entropy Reg



w/ Entropy Reg



Ground Truth

Limitations of the proposed approach

- **Negative alignment:** Feature alignment useful for aligning visually similar features, but classes which are private to a dataset might get wrongly aligned to an unrelated class.
- **Label Prototypes:** Entropy regularization depends largely on the label prototypes, which are unreliable with limited data. Should automatically compute label clusters.

Conclusions and future research

- Presented a method to train joint models which are universally applicable to multiple datasets.
- Entropy regularization helps in within dataset and cross dataset semantic transfer and feature alignment using unsupervised data.
- Future work includes extending the current setting to domain generalization and zero-shot target adaptation.

Thank You!



Any Questions?